

MLMQ

Mat Luke Message Queuing

Advanced Systems Lab – Milestone 1

Matthias Ganz & Lukas Elmer

HS 13/14

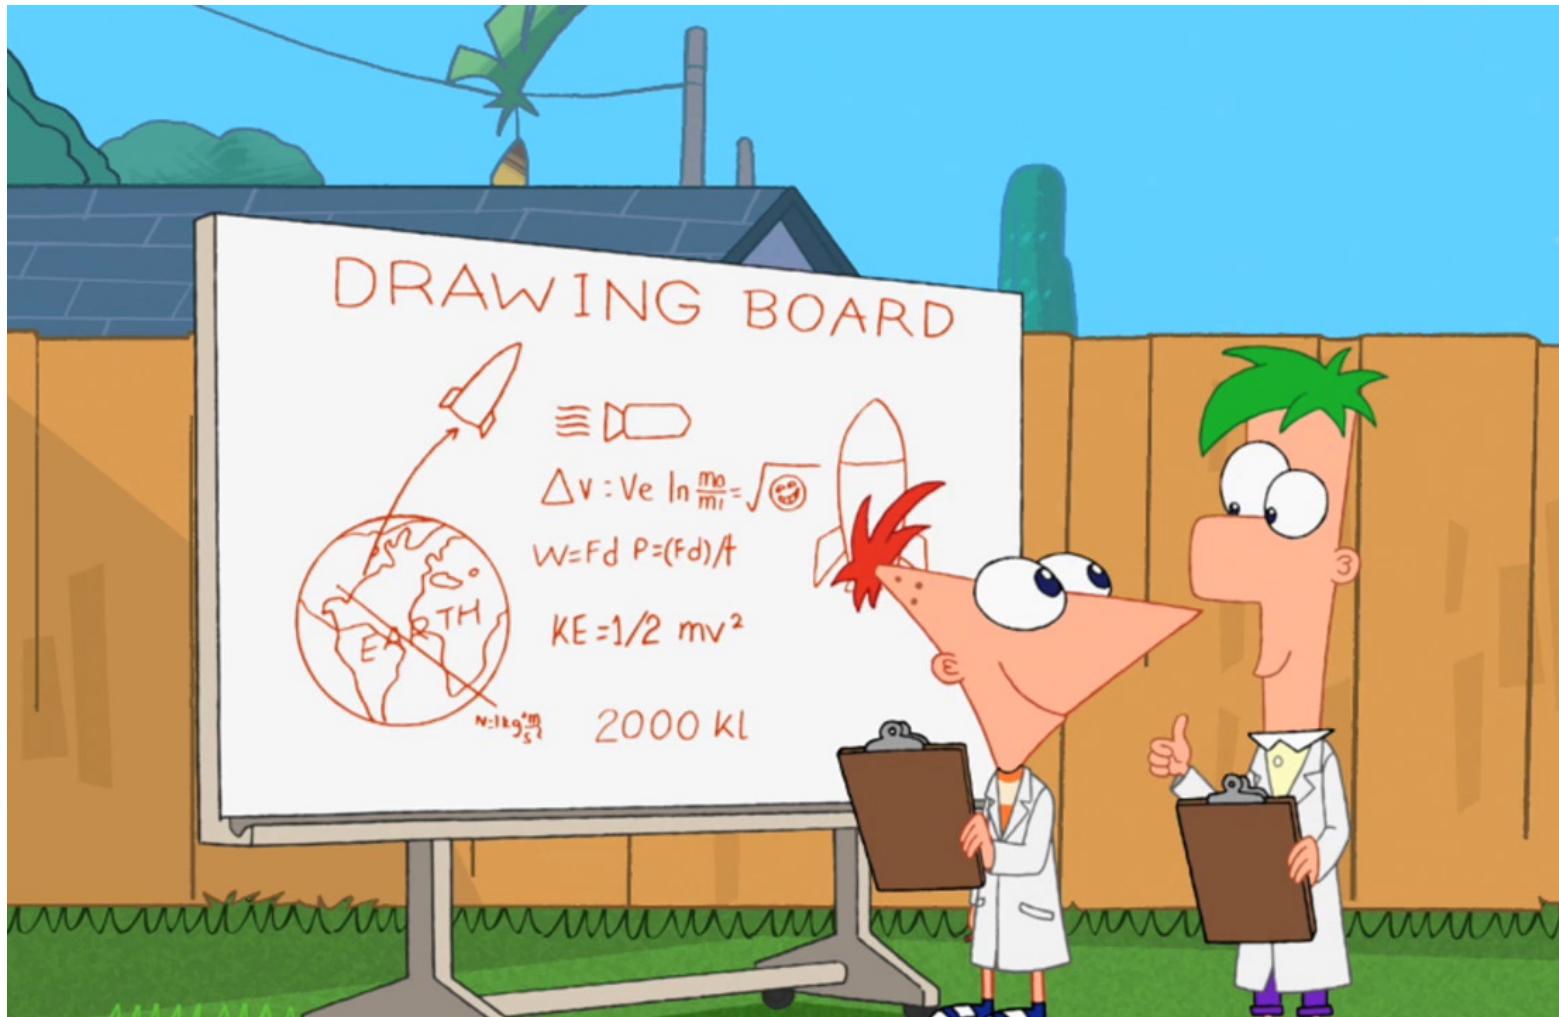
ETHZ



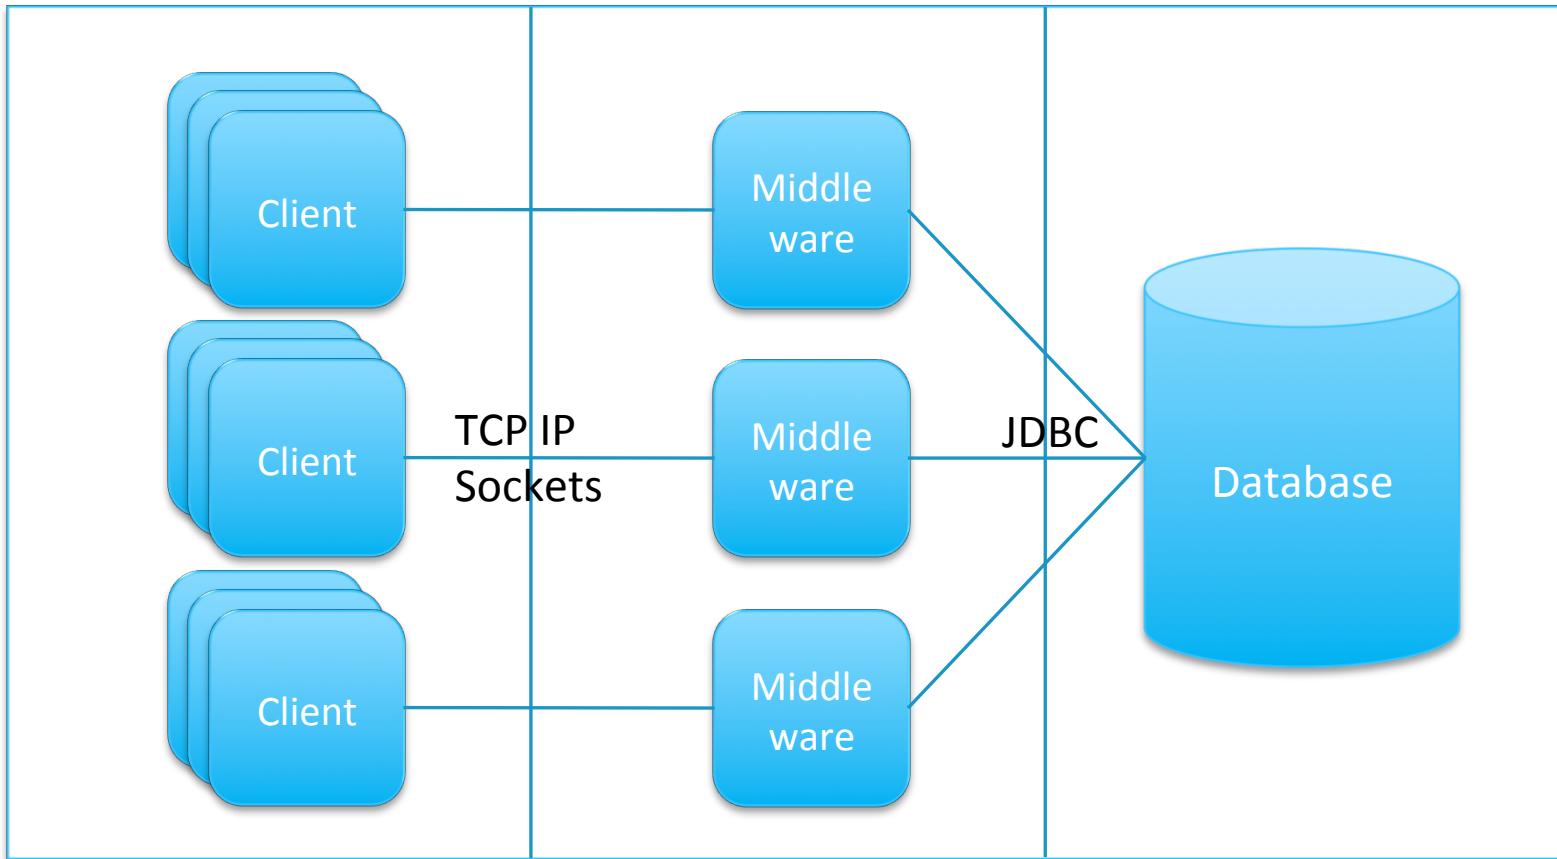
Agenda

- * Design
- * Experiments
- * Lessons Learned about the System

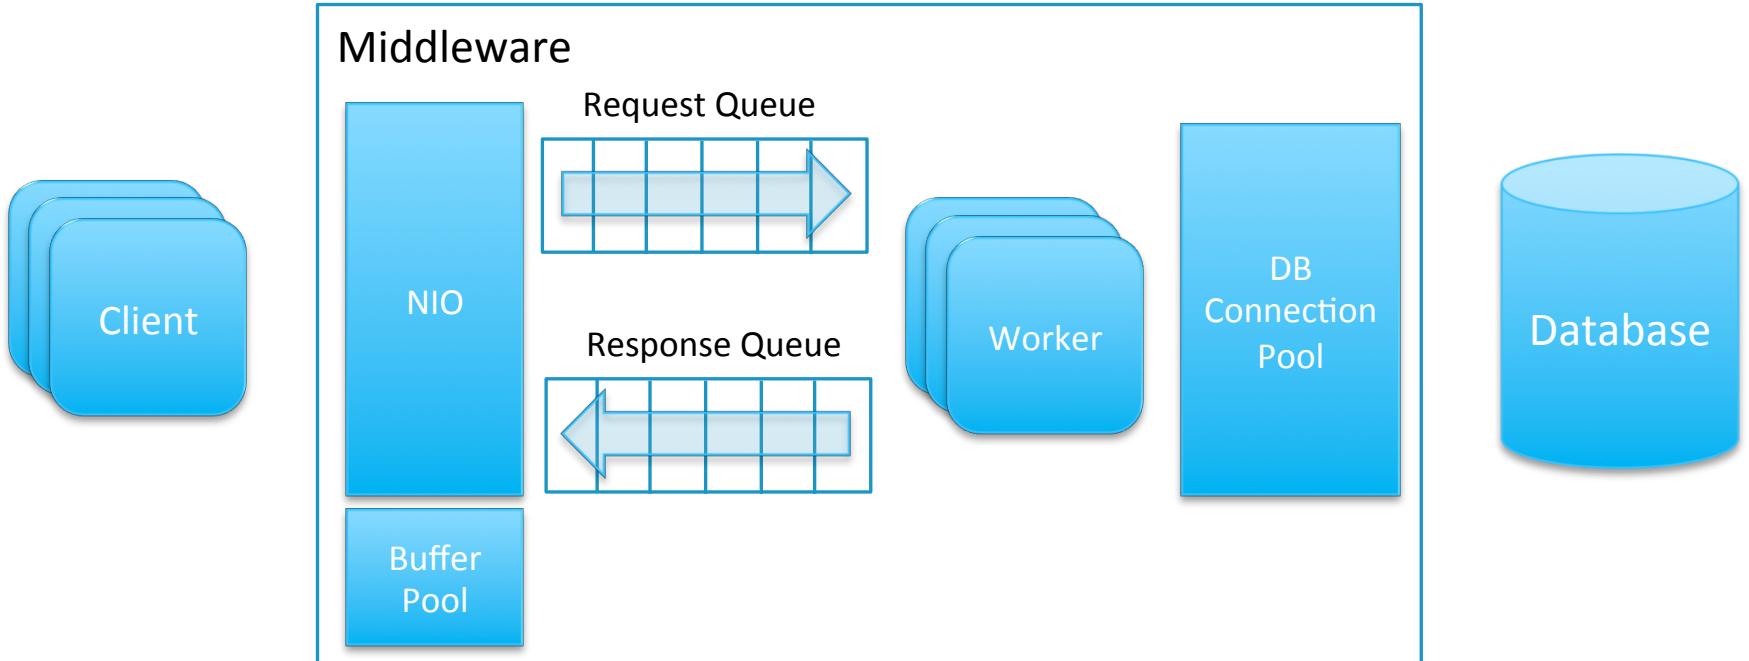
Design



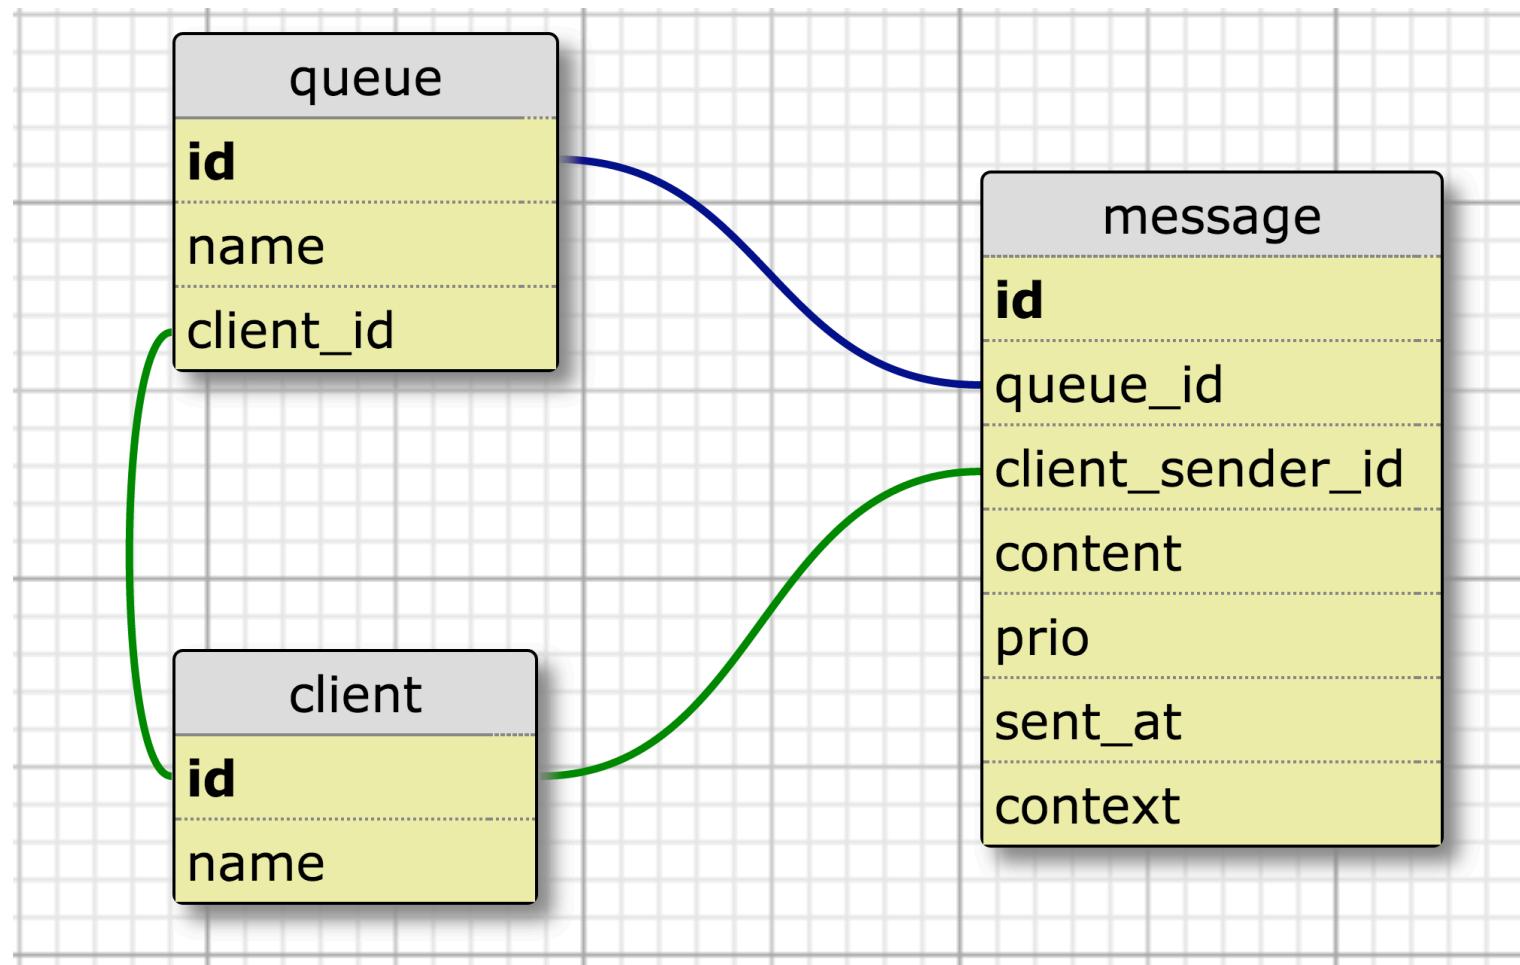
Design – Overview



Design – Middleware



Design – Database



Design – Database Interface

- * Indices

- * Queue: client_id
- * Client: name
- * Message: queue_id, prio, sent_at, client_sender_id, context

- * Important stored procedures

- * peekMessage
- * dequeueMessage
 - * Locks a specific record for dequeuing

- * Prepared statements & auto commit

Design – Client Interface

ClientDto register()

ClientDto lookupClient(String clientName)

QueueDto lookupClientQueue(long clientId)

QueueDto createQueue(String queueName)

QueueDto lookupClientQueue(String queueName)

void deleteQueue(long id)

void sendMessage(long queueId, byte[] content, int prio)

void sendMessage(long[] queueIds, byte[] content, int prio)

void sendMessageToClient(long clientId, byte[] content, int prio)

long sendRequestToClient(long client, byte[] content, int prio)

long sendResponseToClient(long clientId, long context, byte[] content, int prio)

int queuesWithPendingMessages(List<QueueDto> queues, int maxNumQueues)

MessageDto peekMessage(MessageQueryInfoDto messageQueryInfo)

MessageDto dequeueMessage(MessageQueryInfoDto messageQueryInfo)

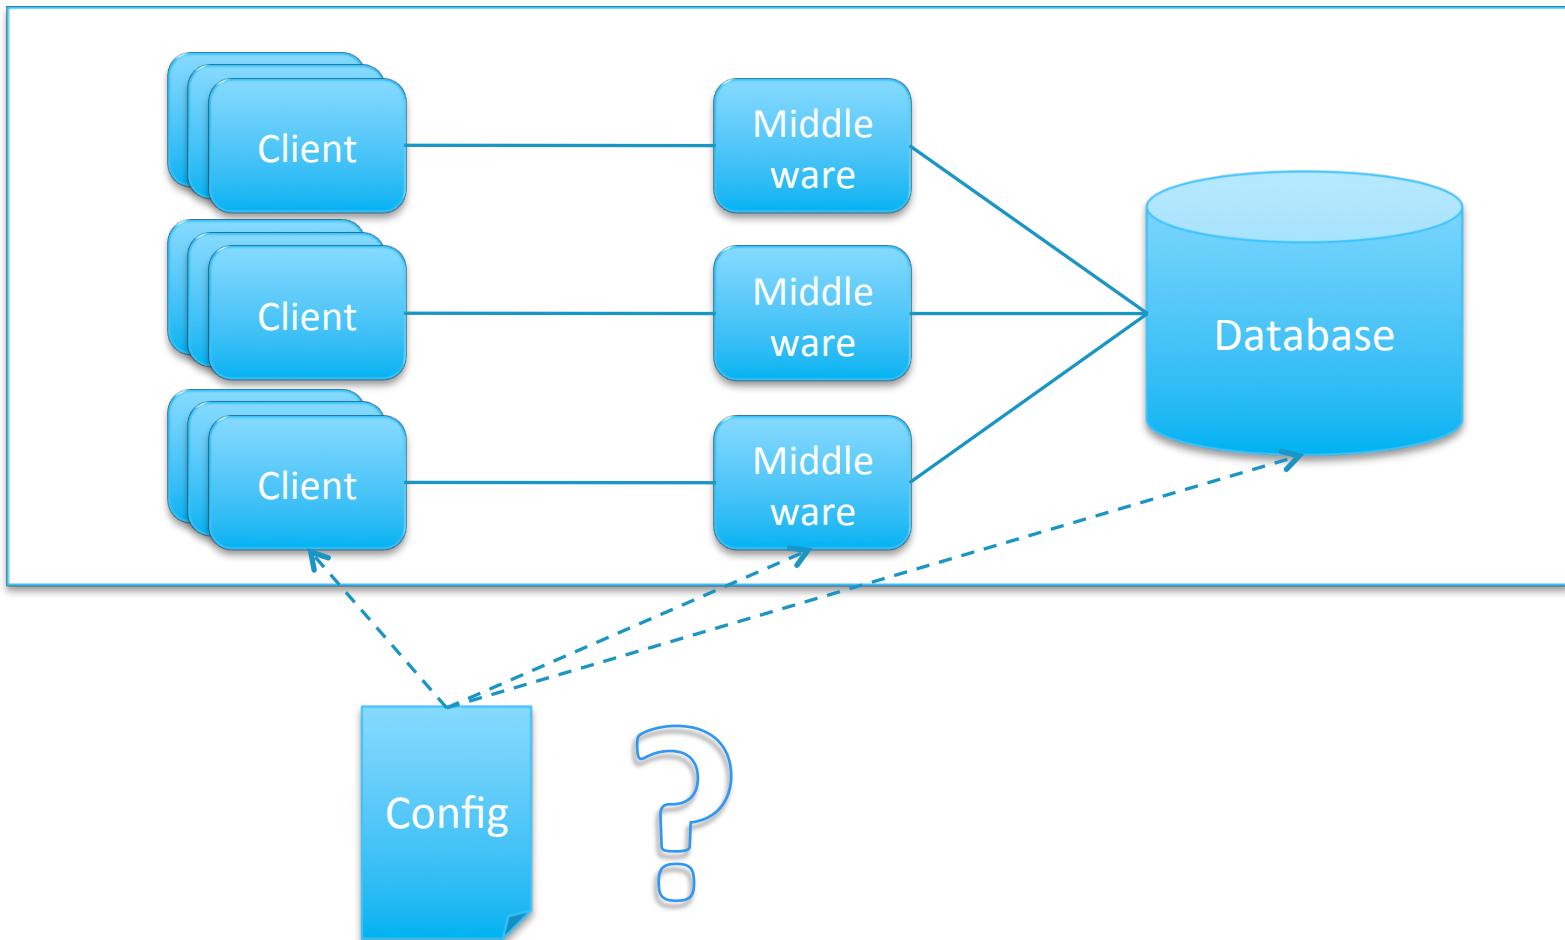
Experiments



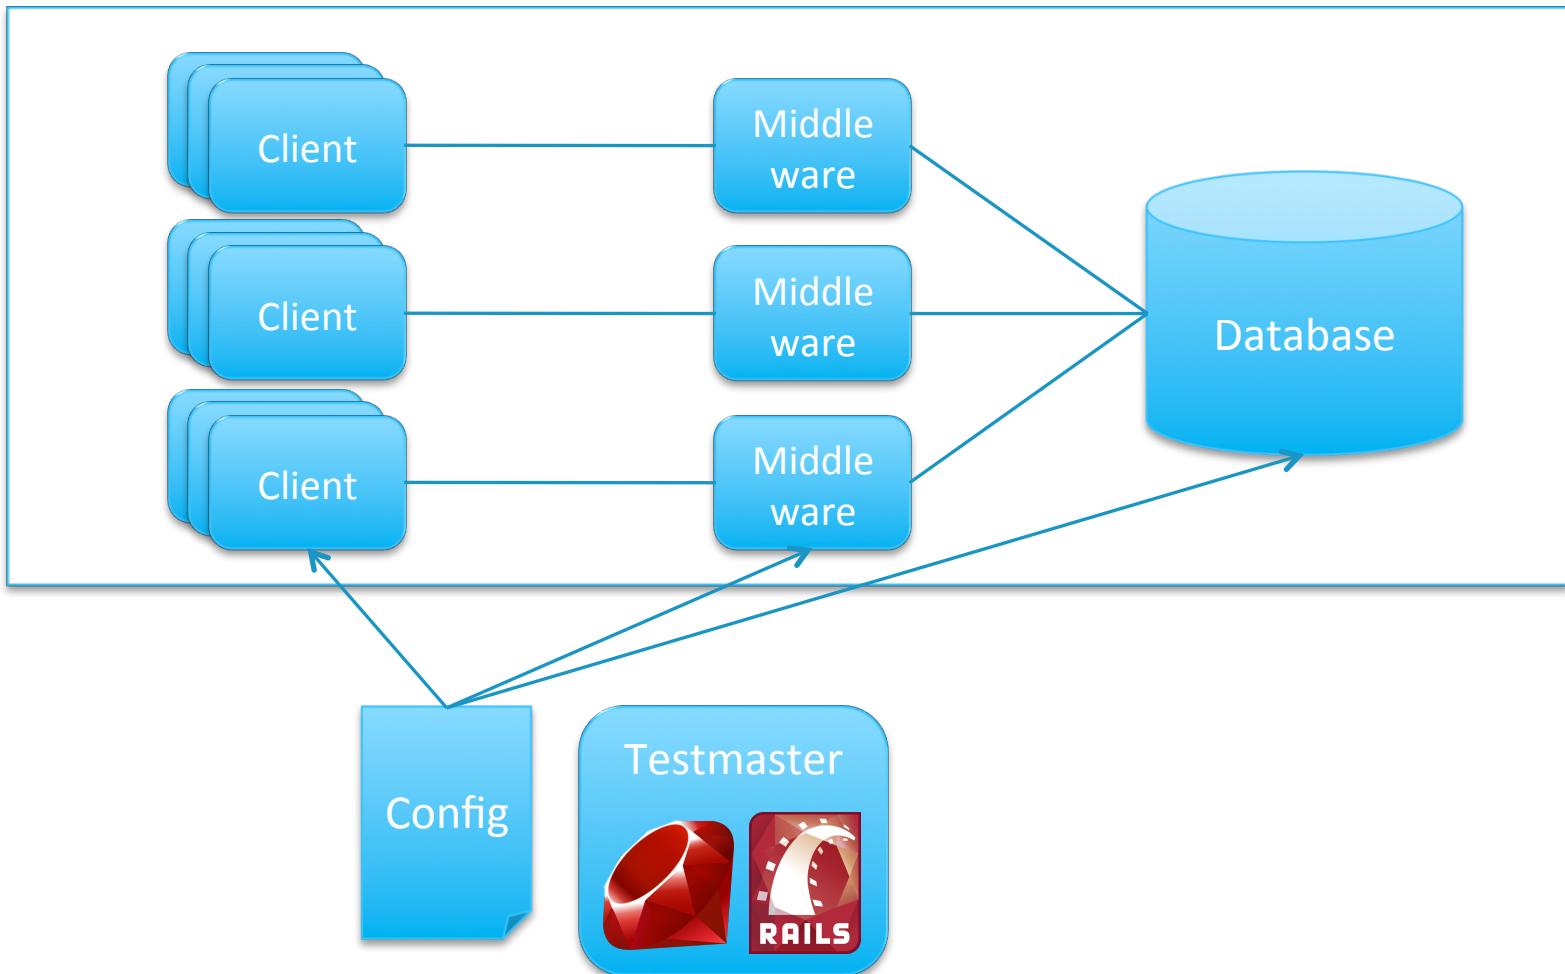
Experiments – Summary

- * Stability
 - * 2h Test
- * Best configuration
 - * Throughput vs. response time
 - * 2^k Test
- * Bottlenecks, optimization Points
 - * Which component spends how much time
- * System limit
 - * Load tests

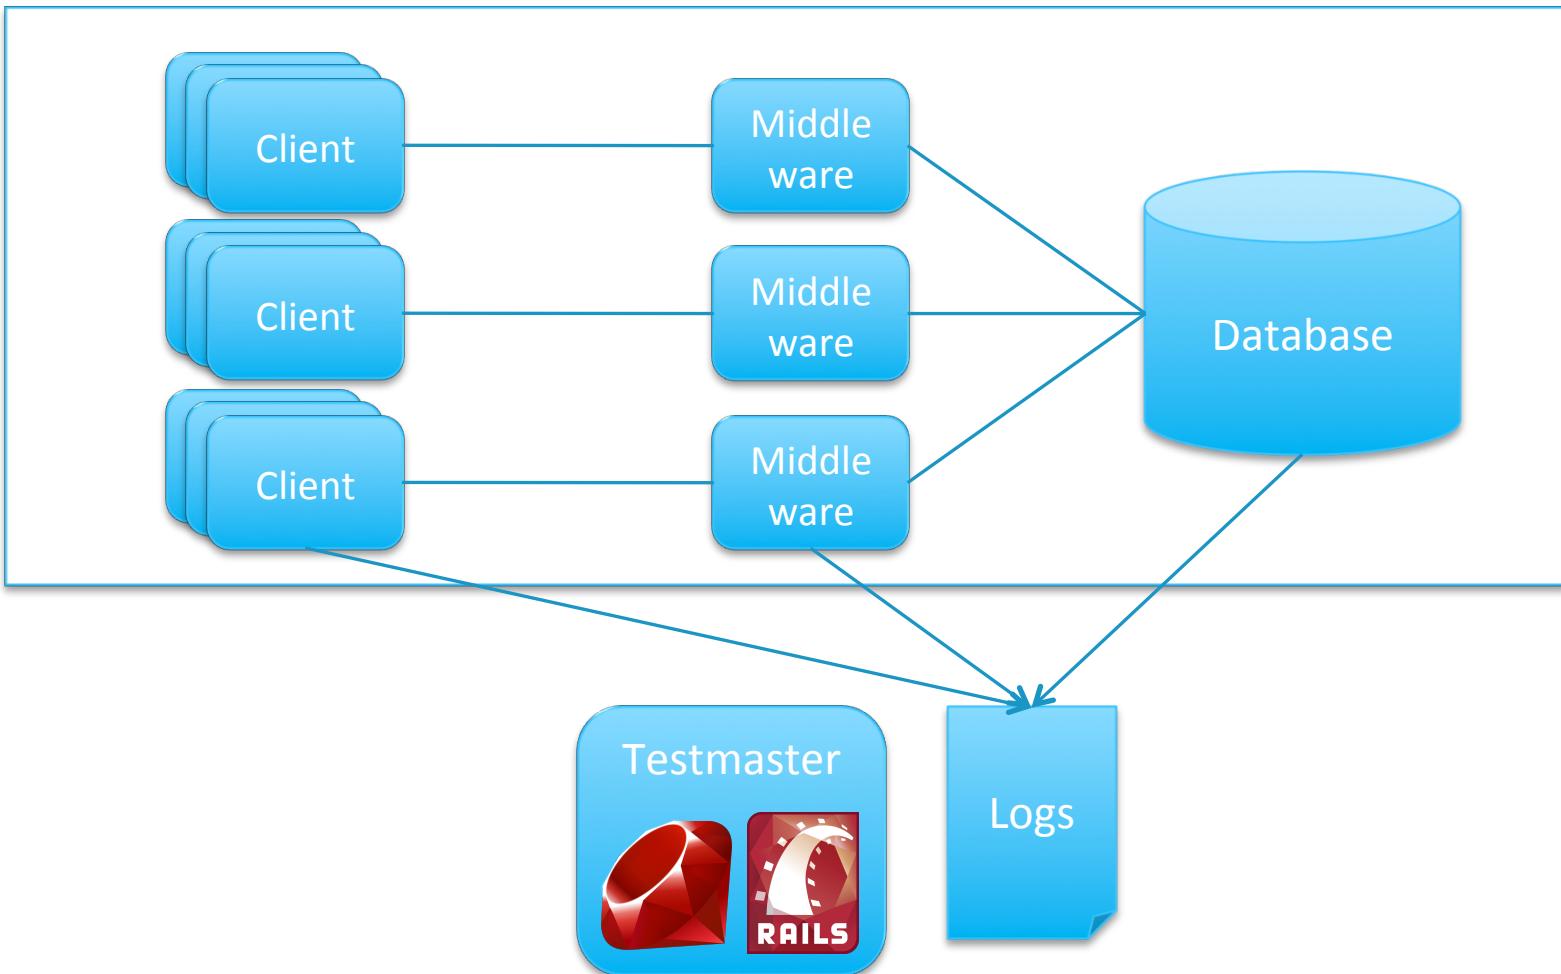
Experiments – Setup



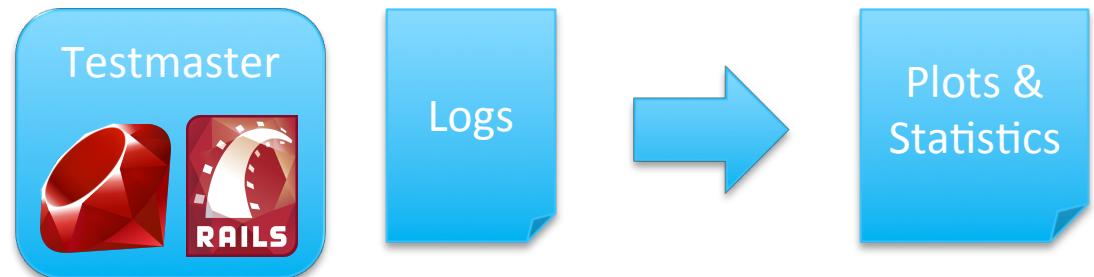
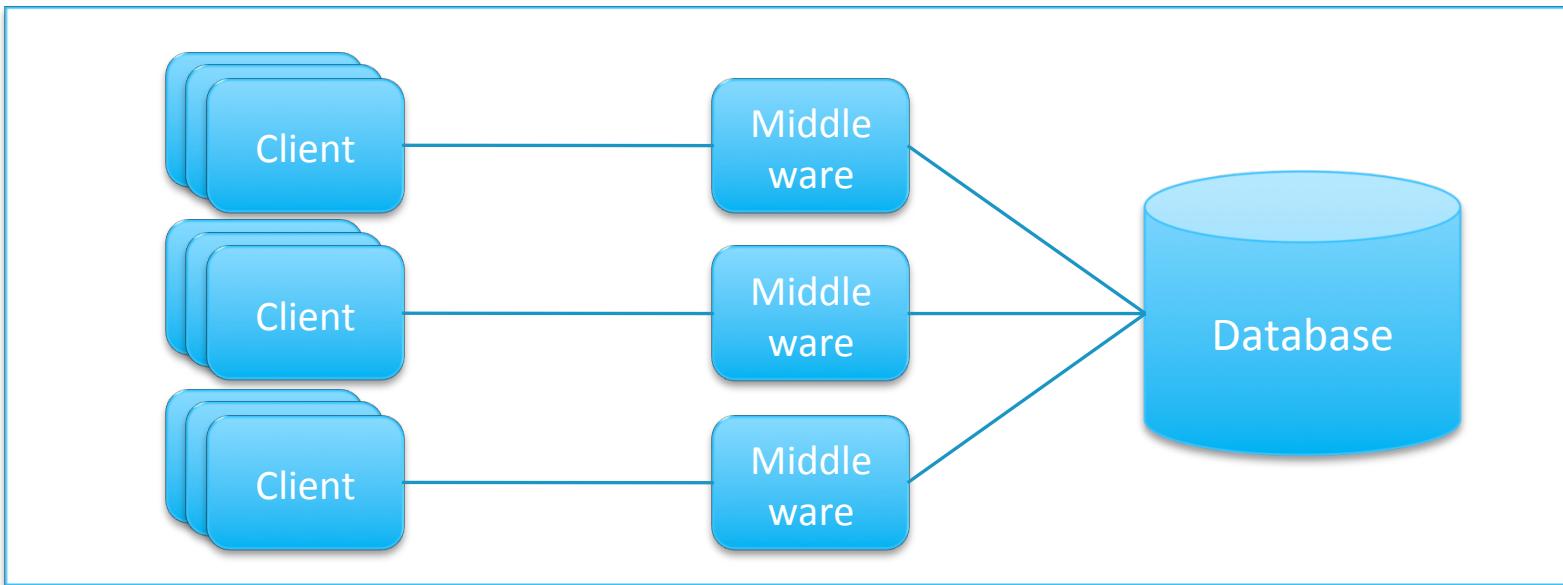
Experiments – Setup



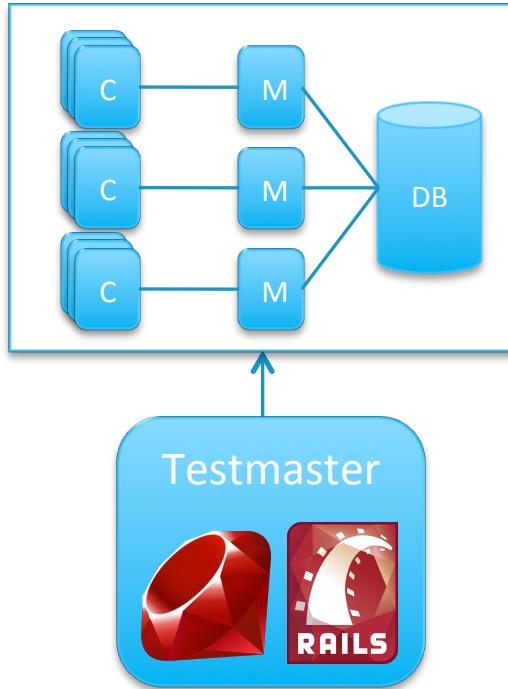
Experiments – Setup



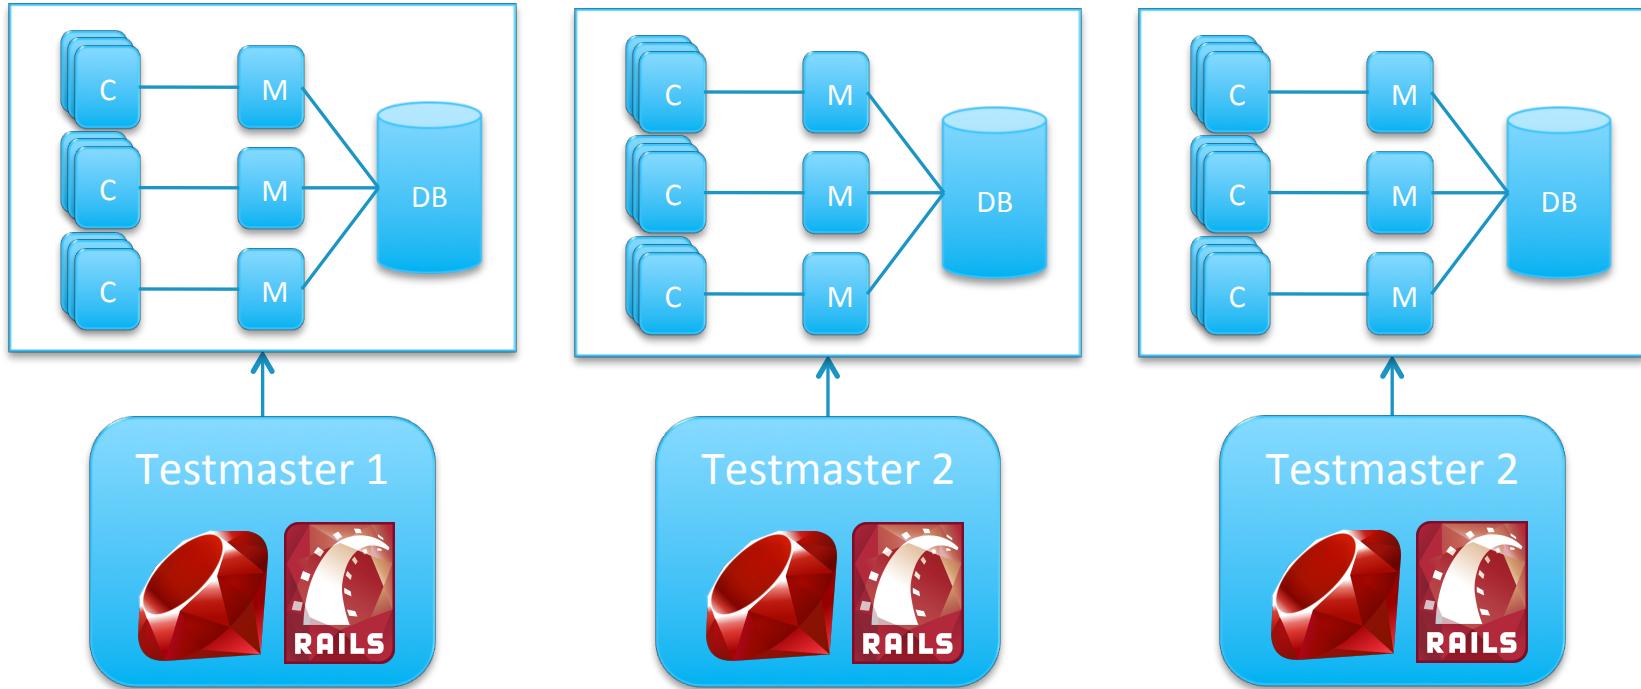
Experiments – Setup



Experiments – Sequential Test



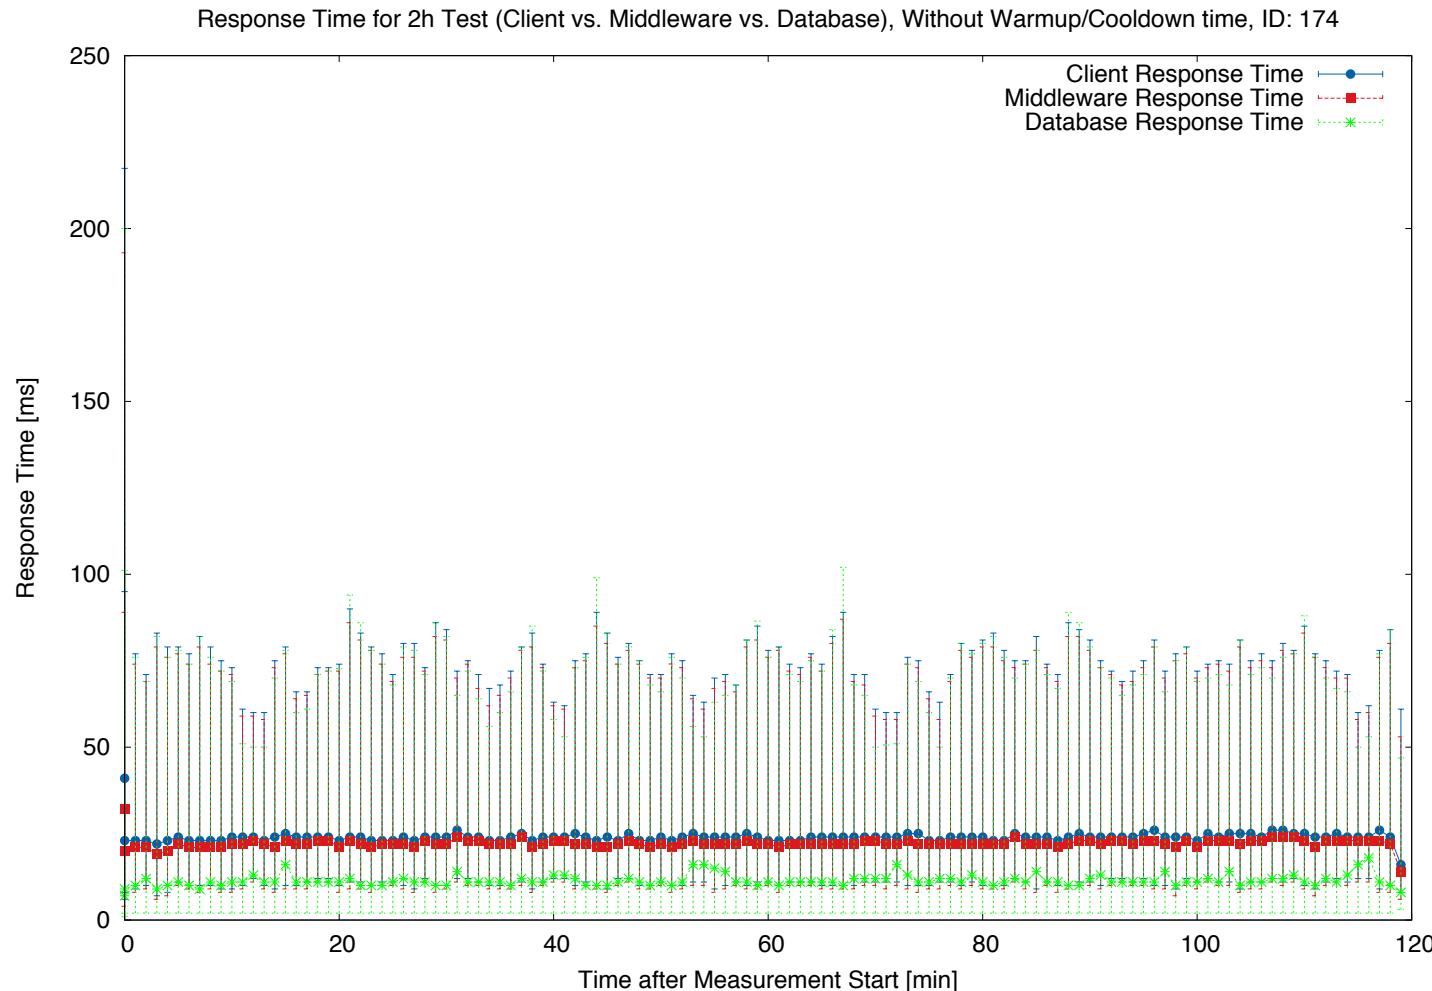
Experiments – Parallel Tests



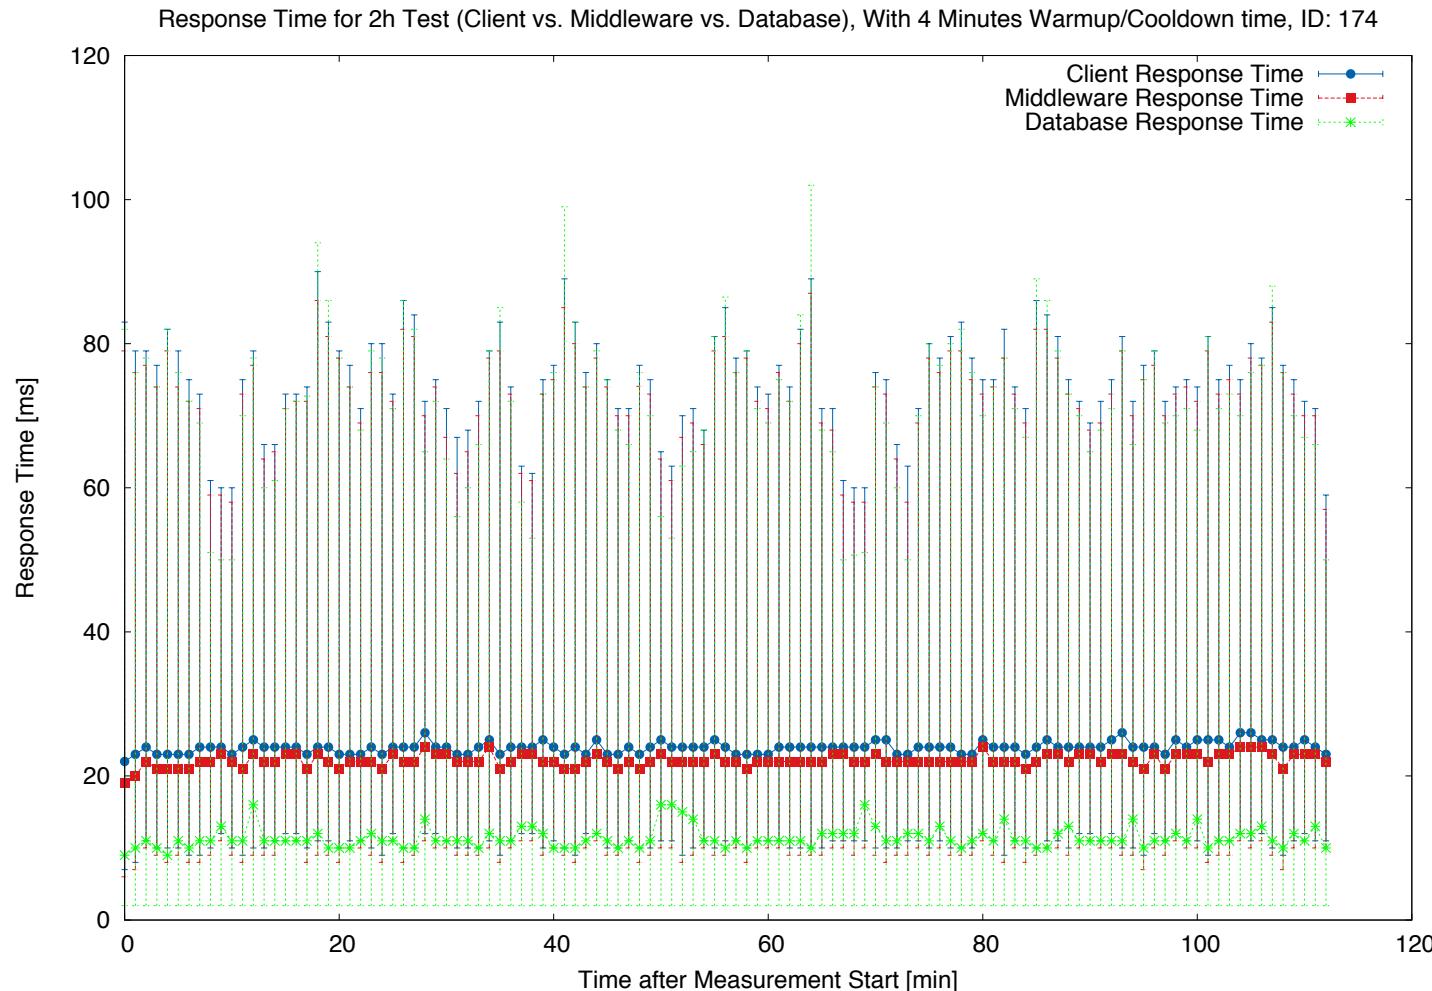
Cannot to 20 Amazon EC2 instances

→ Request to Increase Amazon EC2 Instance Limit

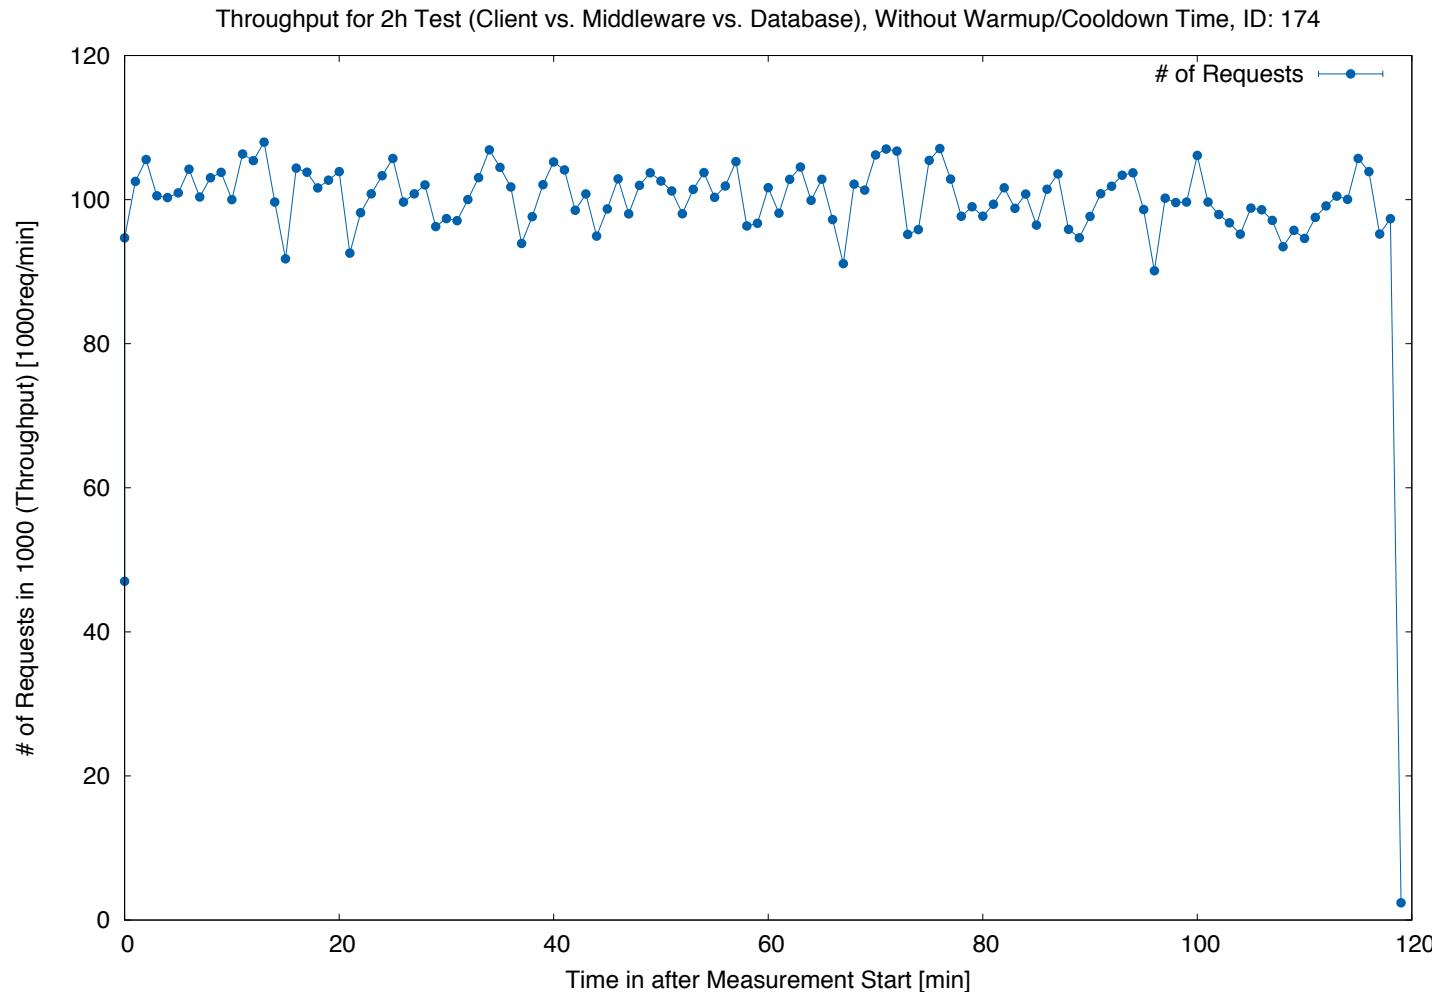
Experiments – 2h Test without warmup / cooldown time, response time



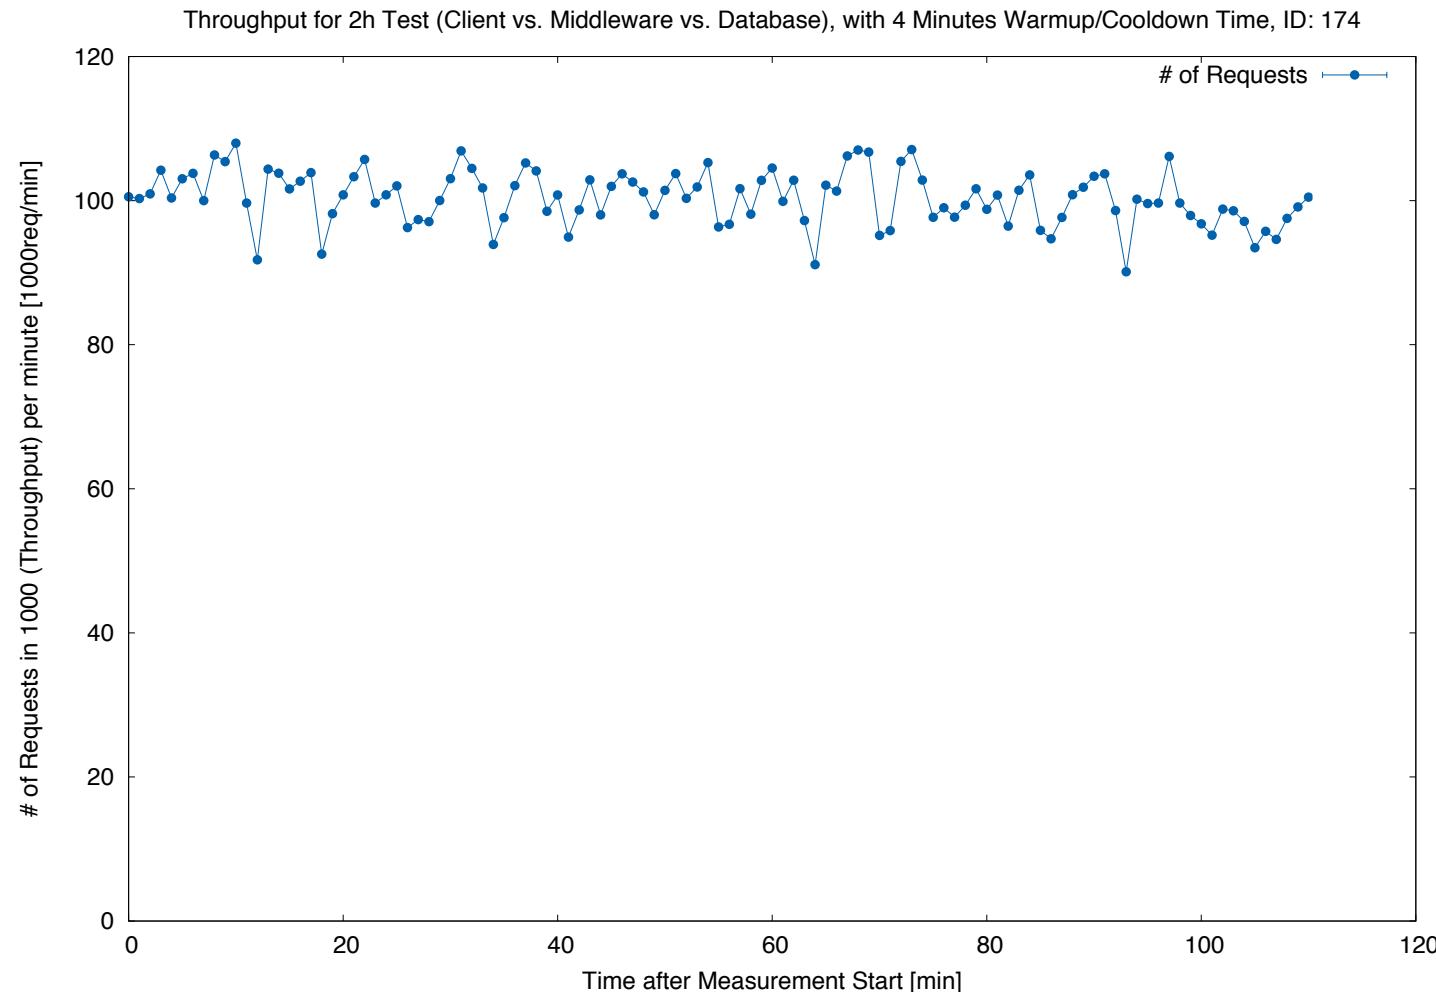
Experiments – 2h Test with 4 minutes warmup / cooldown time, response time



Experiments – 2h Test without warmup / cooldown time, throughput



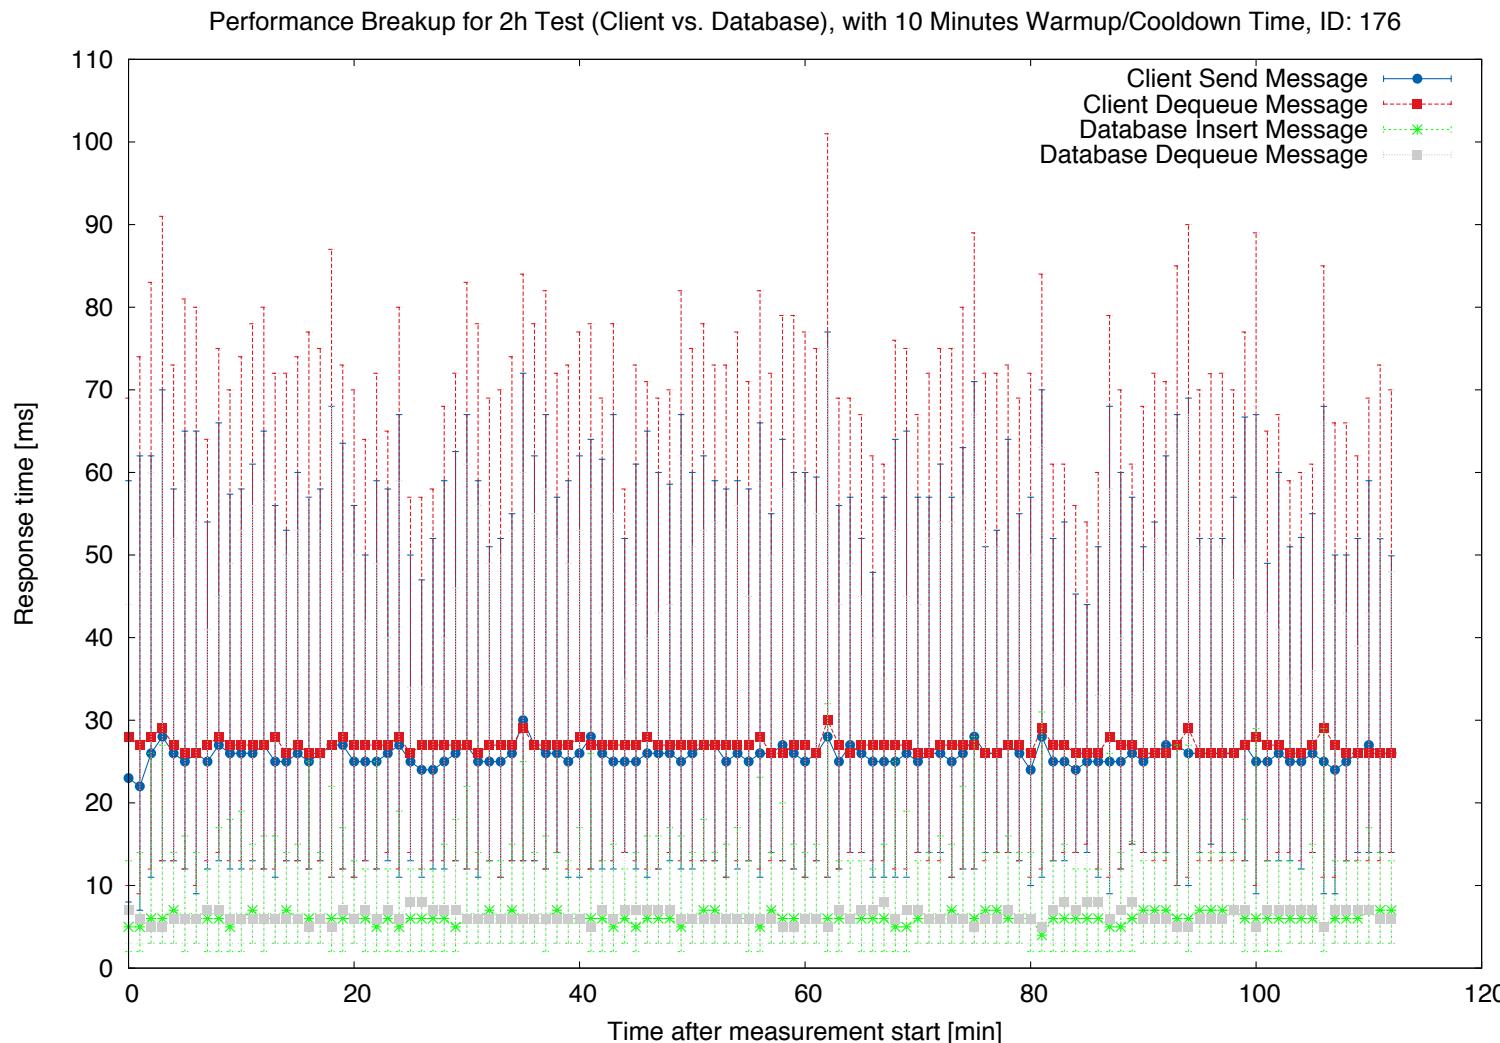
Experiments – 2h Test with 4 minutes warmup / cooldown time, throughput



Experiments – 2h Test

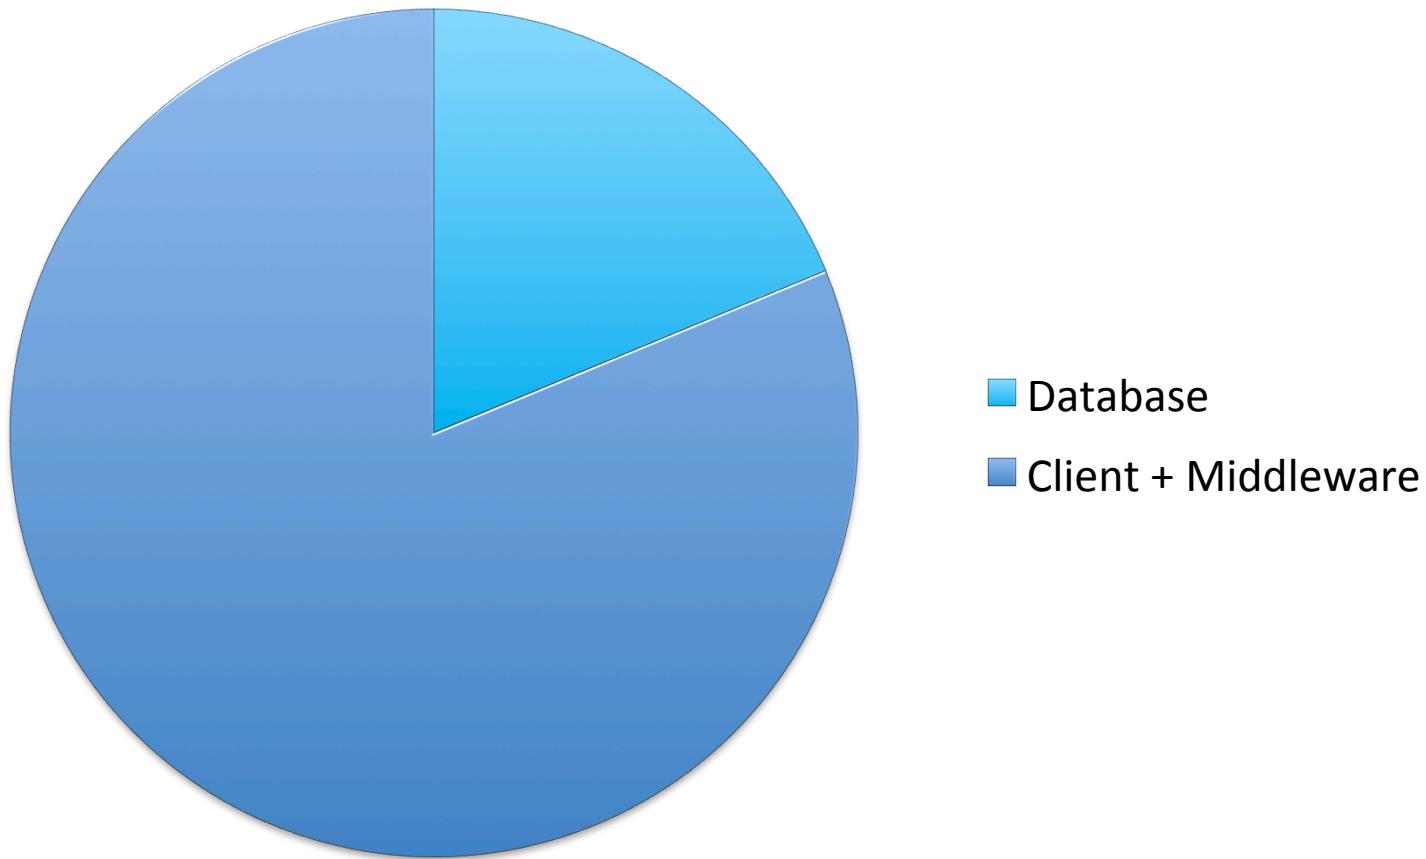
- * For the 2h Test configuration
 - * For 95% of all **sendMessage** requests, the response time will be **under 50 ms**
 - * For 95% of all **dequeueMessage** requests, the response time will be **under 62 ms**

Experiments – Results, 2h Test



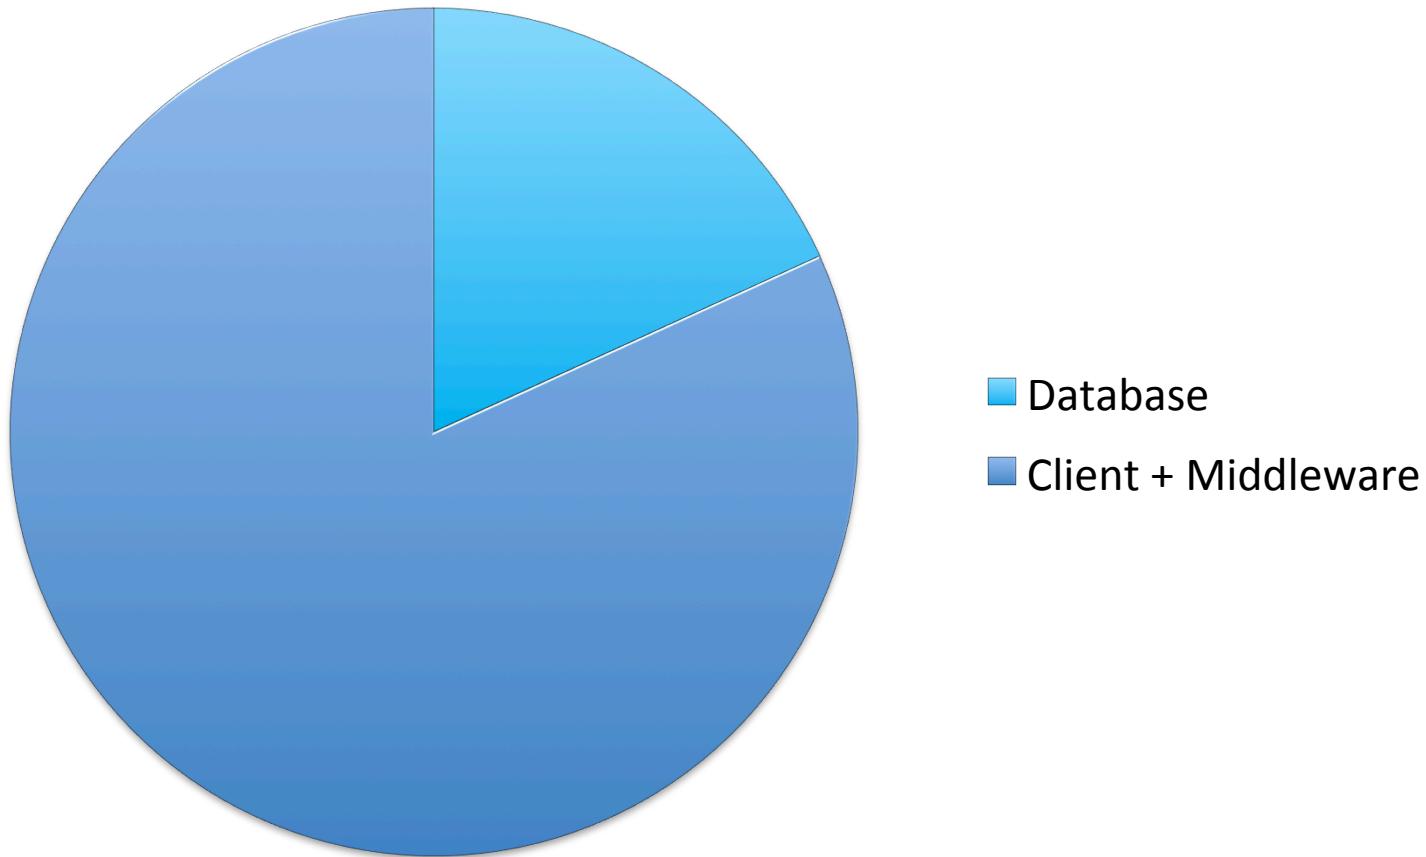
Experiments – Results, 2h Test

Time spent for sendMessage (2h Test)

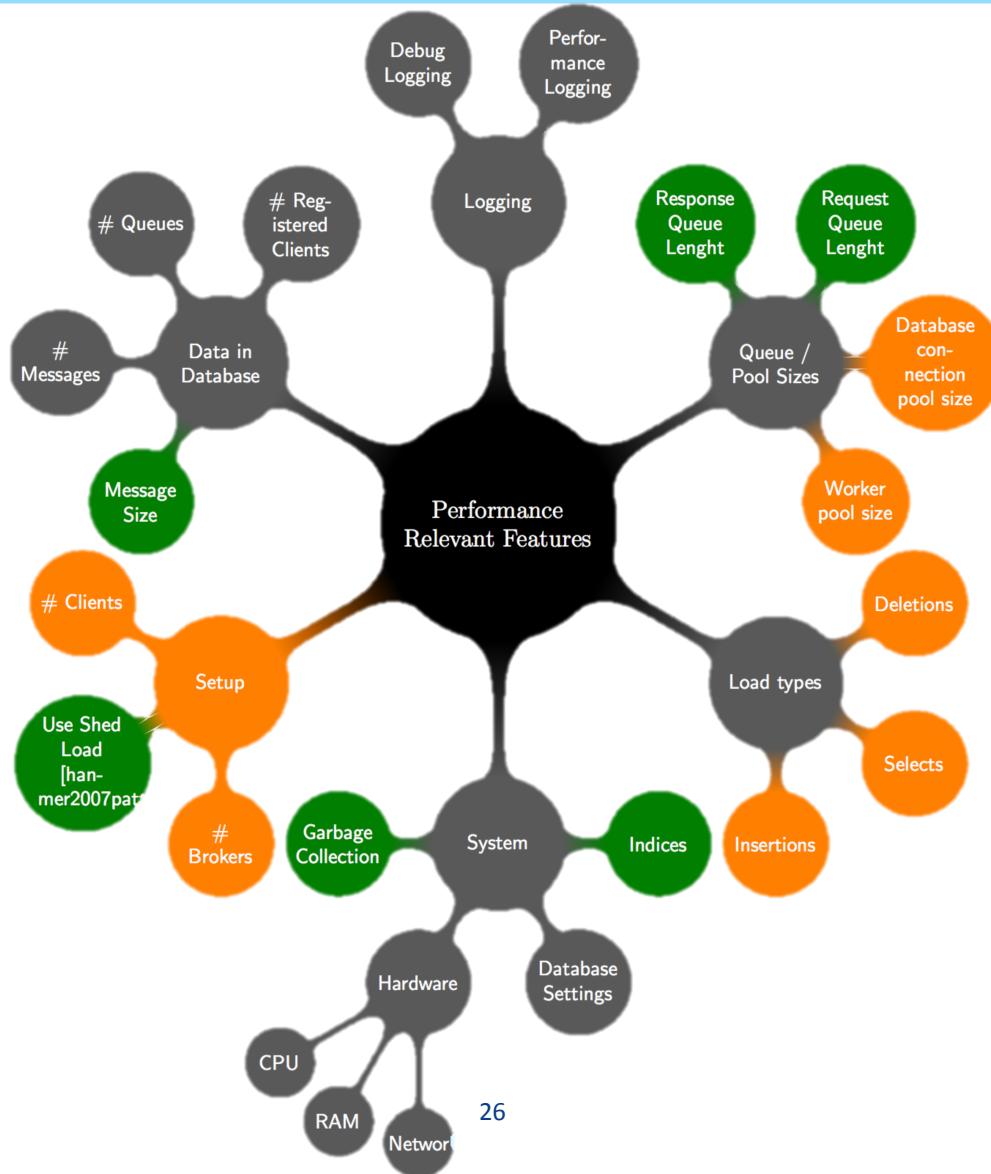


Experiments – Results, 2h Test

Time spent for dequeueMessage (2h Test)

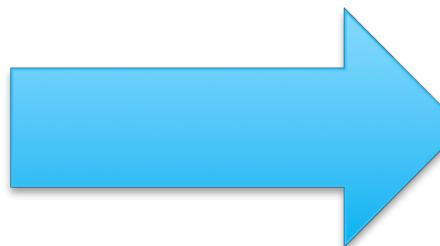


Experiments – Factors



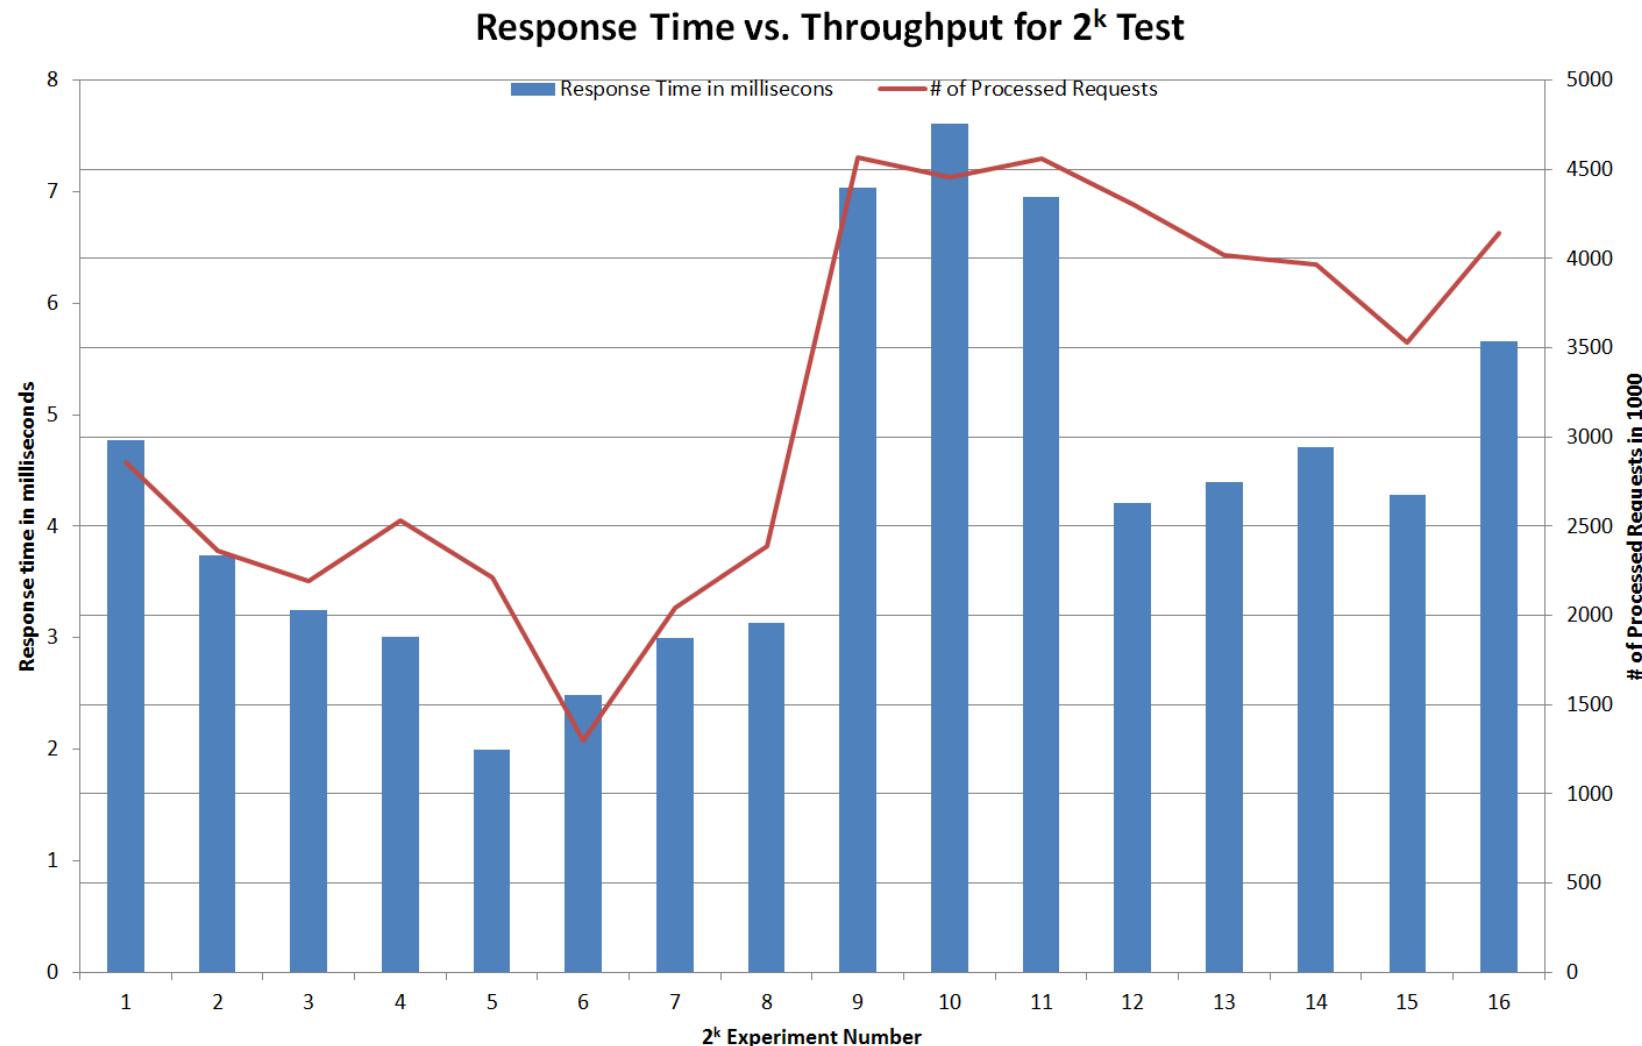
Experiments – Primary Factors

- * # Clients
- * # Brokers
- * DB Connection Pool Size
- * Worker Pool Size



2^k Test

Experiments – Results: 2^k Test

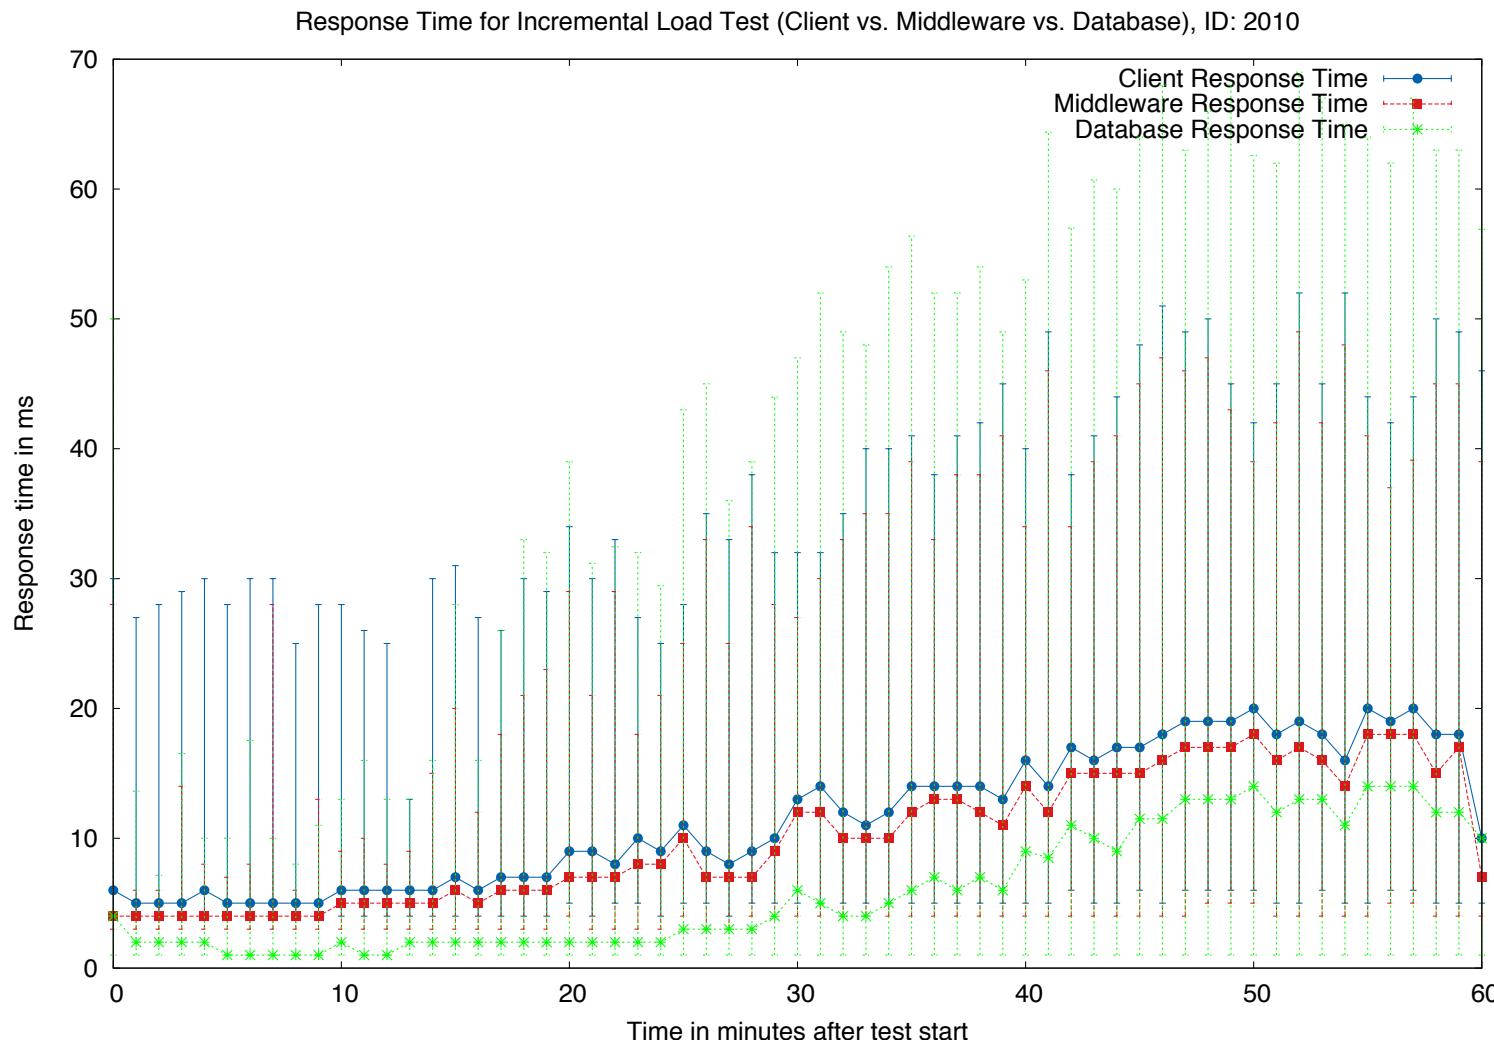


Experiments – Results: 2^k Test

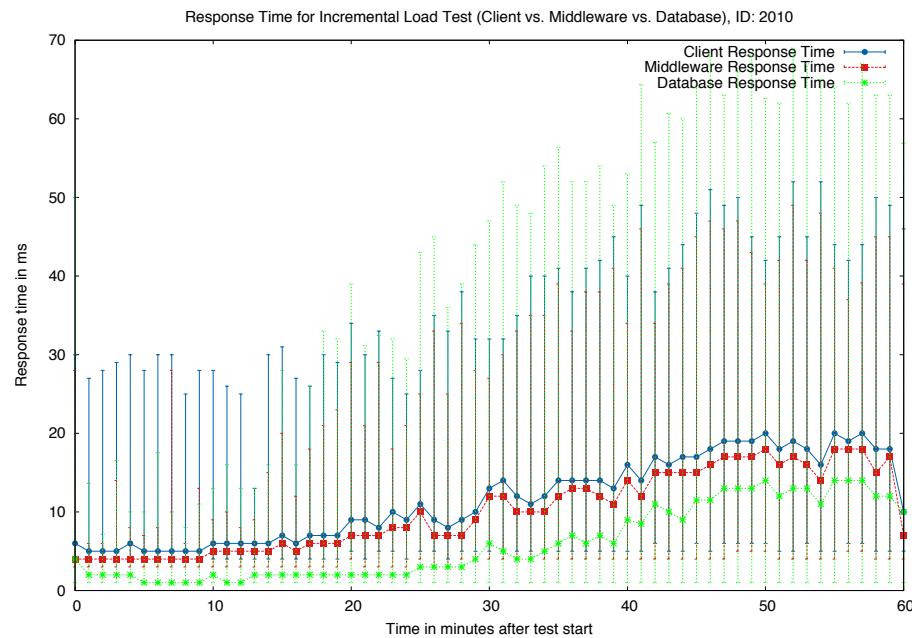
* Interpretation

- * Higher throughput → higher response time

Experiments – Incremental Load Test, Response Time

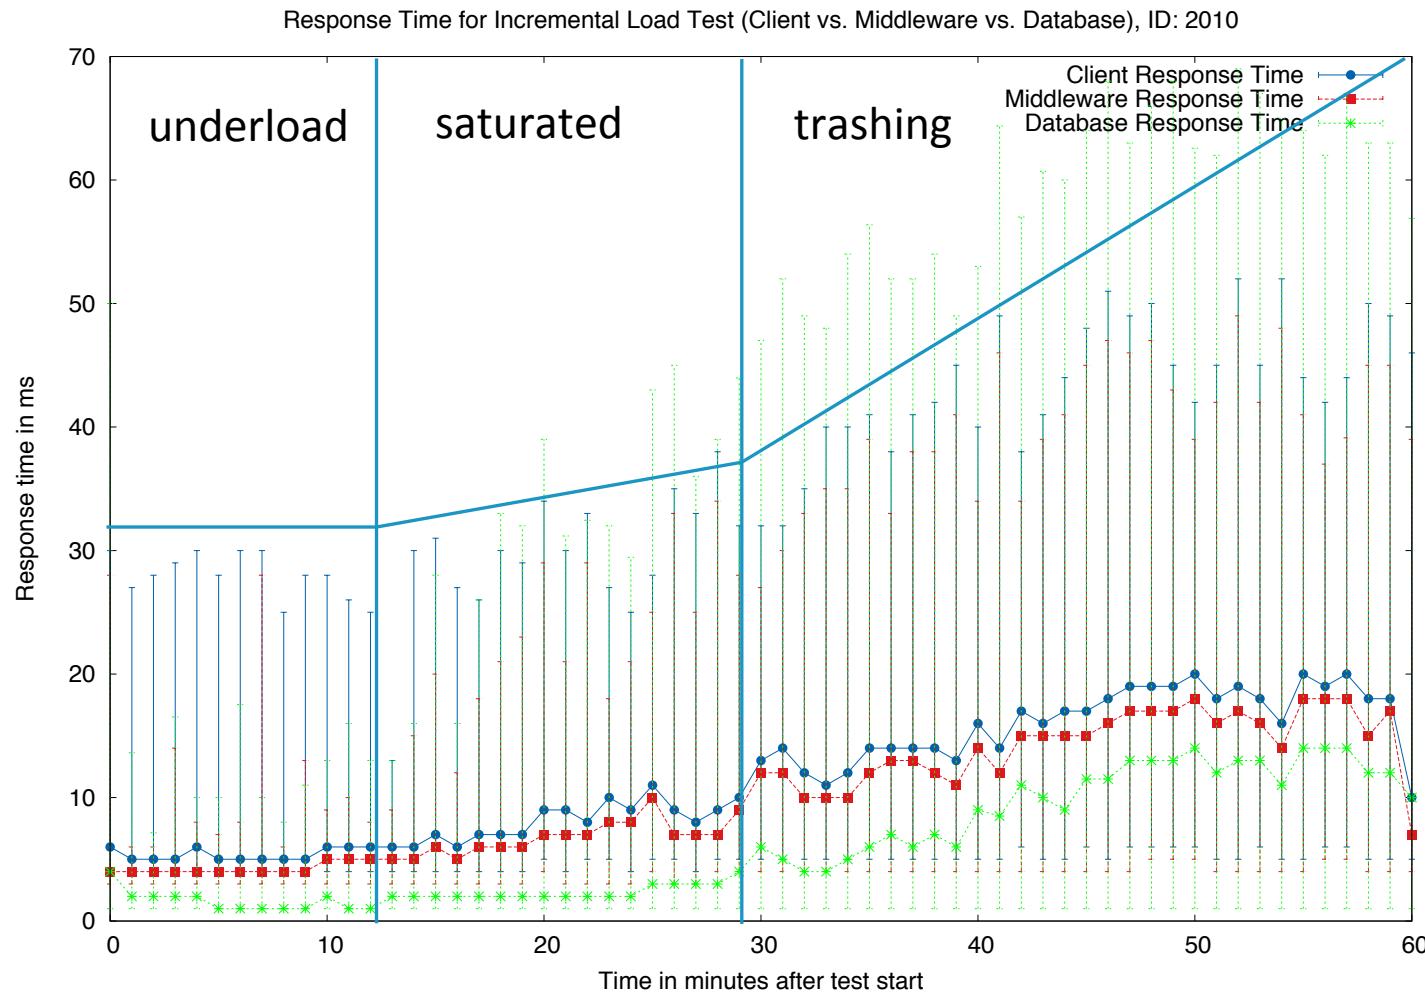


Experiments – Incremental Load Test, Response Time

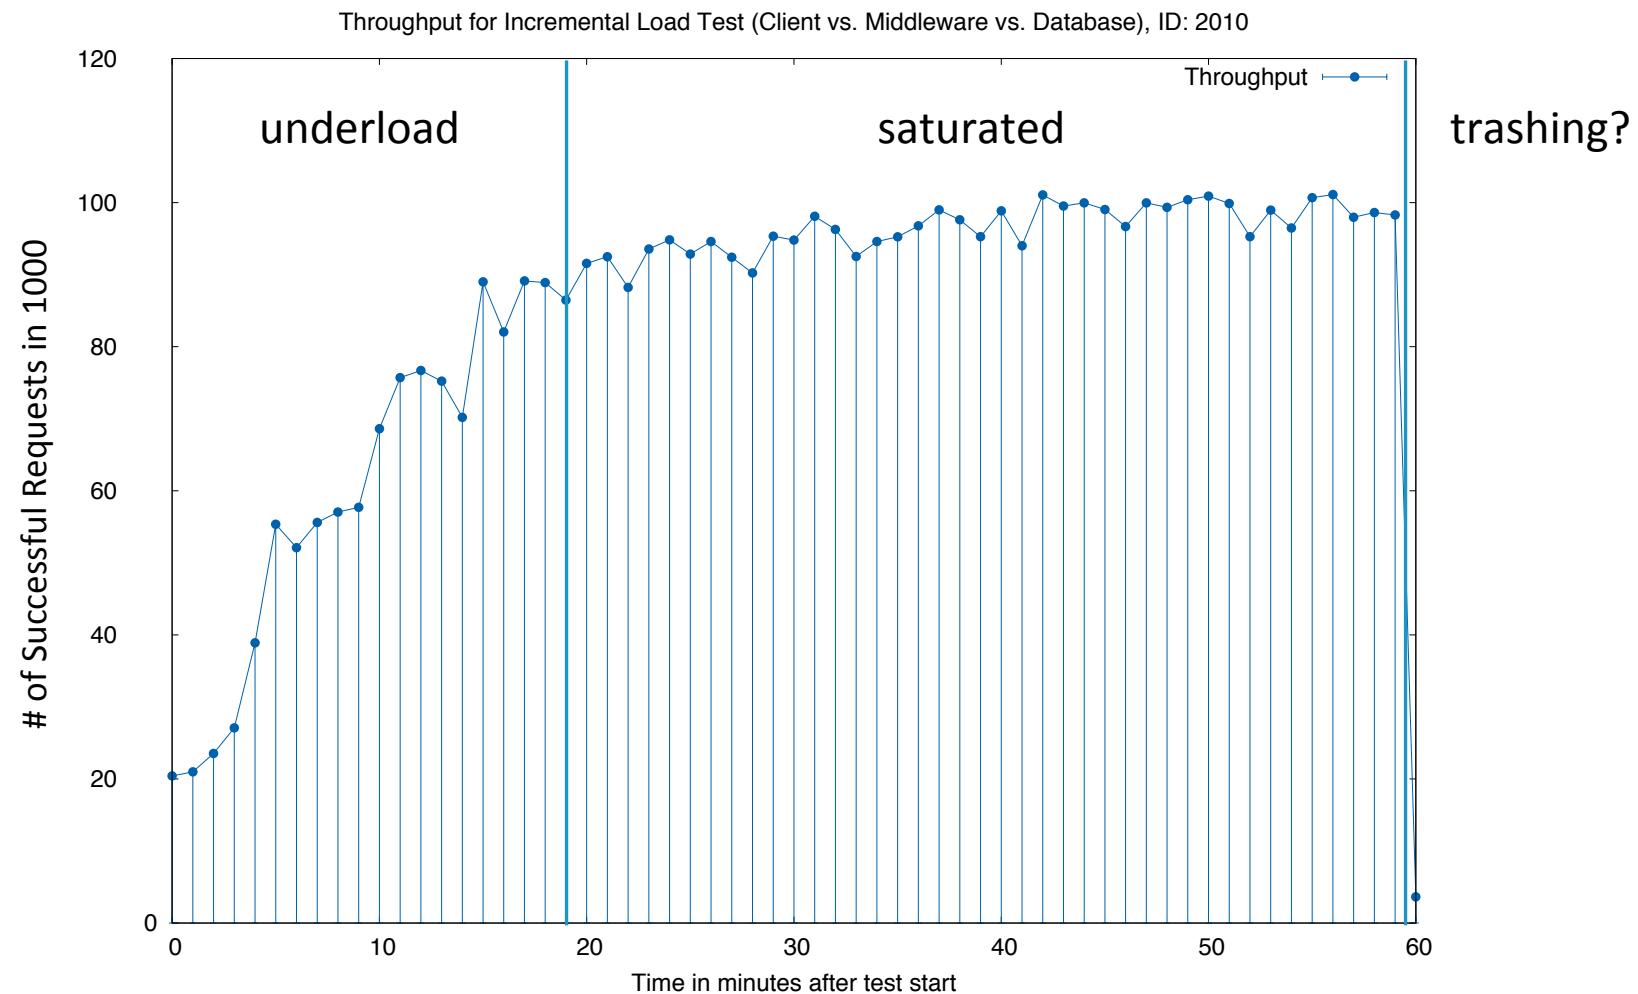


- * Interpretation: Higher load →
- * Slower database
- * Higher database response time variance

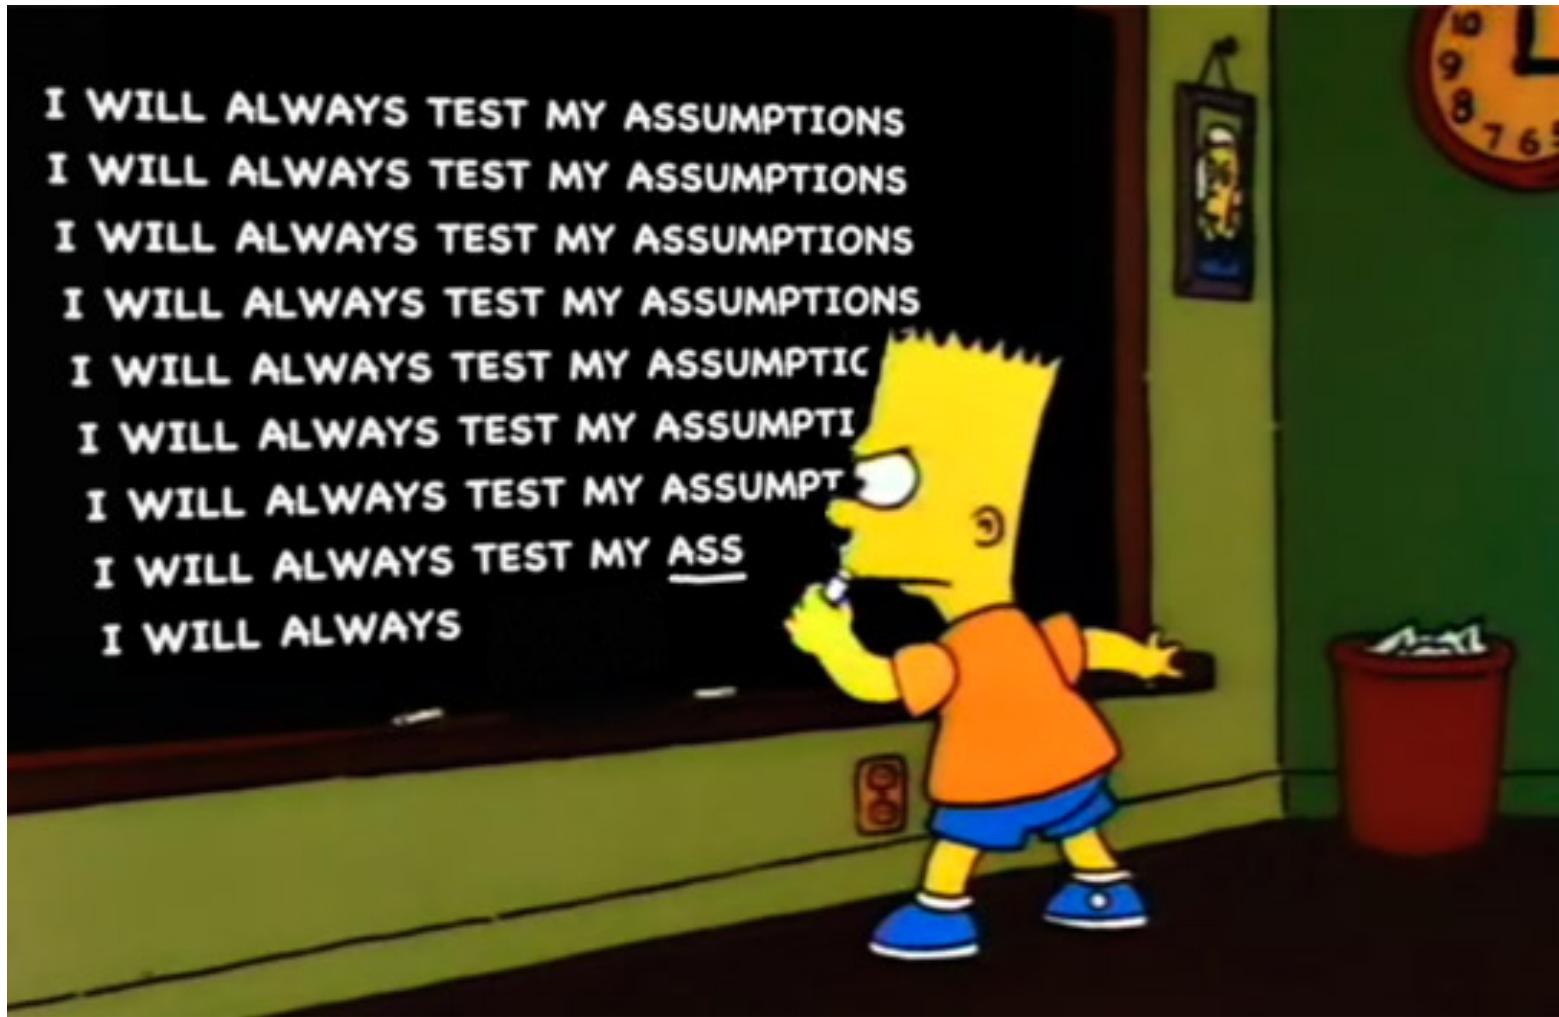
Experiments – Incremental Load Test, Response Time



Experiments – Incremental Load Test, Throughput



Lessons Learned about the System



Lessons Learned about the System

- * Improvements during this milestone

- * Response time stability
 - * Throughput

- * Saturated system

- * => higher throughput
 - * => higher variance in response time

- * Potential bottleneck

- * Database

Thank you

Questions